


H17: Behavior is a Markov state model with a few metastable sets; its mixing time is a per-period order parameter

Vivian Rogers, with Claude Opus 5.5 · agent-swarm physics · Updated 2026-10-03

Complete 40% Faithfulness 1/5 Usefulness 1.5/5


Bottom line. Behavior on action classes is a non-Markov mix of tool modes and idling. Its slowest timescale is a stable but uninformative period statistic (not an order parameter for stuckness), and the content-level test awaits the Jev states.

level: hypothesis · pipeline checks 100% · updated 2026-10-03



WHAT HAPPENED

Is agent behavior a Markov state model (MSM) with a few slow metastable sets, and is its slowest timescale t_2^* a per-period order parameter, slowest in stuck periods? Round 1 (exploratory) fitted MSMs on 6 action classes per active minute in 27 non-holdout periods (17 regime I, 2 regime II, 8 regime III). Synthetic tests showed the estimators recover t_2 and planted sets at village sample sizes. On real data: CK fails in 27/27, and $t_2(\tau)$ rises $\times 4-7$ from 1 to 15 min, so the chain is not Markov. t_2^* (5–24 min) beats a sojourn-preserving null by only $\times 1.06-1.62$, which passes in 3/8 regime-III periods. Pooling agents inflates t_2^* by $\times 1.28$. t_2^* does not track errors or output, regime III is *slower* than regime I (reversed), and a “don’t do nothing” prompt (NE07) lengthened it. Verdicts: 13 mixed, 14 failed. Holdout (#32, #45) not run.



Observables. Slowest implied timescale per goal period at $\tau_c = 5$ min. Filled: pooled MSM (95% agent-day CI). Bars: sojourn-null median (keeps each state’s dwell times, shuffles their order). Open: median per-agent t_2 . The pooled MSM beats the sojourn null in every period, but only modestly, and it sits above the typical agent (agent-mixture inflation). Regime III (red) is slower than regime I (blue). #37’s peak is agents idling before their first and after their last action of the day.

PHYSICAL MODEL

Kinetic Potts / non-equilibrium Markov chain (models 02, 10). State $x_t \in \{\text{browse, type, shell, chat, idle, consolidate}\}$ per agent-minute (H14 scheme); $C(\tau)$: lagged pair counts within agent-days; $T(\tau)$: non-reversible MLE; t_k : implied timescales from eigenvalues λ_k ; χ : PCCA+ memberships (Röblitz-Weber) of m metastable sets; $q^\pm$: forward/backward committers, $\delta = 0$ under detailed balance; p_t : soft (Jev) posteriors.

$$T_{ij}(\tau) = C_{ij}(\tau) / \sum_k C_{ik}(\tau), \quad t_k(\tau) = -\tau / \ln |\lambda_k(T(\tau))|$$

$$t_2^* \equiv t_2(\tau_c = 5 \text{ min})$$

$$\text{CK: } P_{A \rightarrow A}(k\tau) \stackrel{?}{=} [\pi_A^\top T(\tau)^k]_A$$

$$T_{cg} = (\chi^\top \Pi \chi)^{-1} \chi^\top \Pi T \chi, \quad \delta_i = q_i^+ - (1 - q_i^-)$$

$$\lambda_k^\tau = \text{eig}[C(1)^{-1} C(1+\tau)], \quad C(\tau) = \sum_t p_t p_{t+\tau}^\top$$

SO WHAT? PRACTICAL VALUE FOR UNDERSTANDING AGENT SWARMS

Candidly: on action classes, this is a curiosity with one methodological payoff. The “metastable sets” are mostly tool modality (GUI vs. shell) and idling, which any log viewer could tell you. Part of the GUI-vs-shell split is GUI-heavy vs. shell-heavy agents. t_2^* is a stable, cheap period statistic, but it does not detect stuck periods (no link to error share or output), so an operator gains no monitor. The useful part is negative and transferable: minute-scale agent behavior is strongly non-Markov, pooled MSMs over heterogeneous agents manufacture slow modes, and LLM-labelled states with $\kappa \approx 0.5$ would bias timescales low by 4–14 $\times$ unless a noise-cancelling estimator is used. Whether content-level (Jev) states show real “stuck” attractors is still open.

H17-behavior-metastable-sets · generated 2026-10-03 19:00 by infra/summaries/build_summaries.py · card hypotheses/H17-behavior-metastable-sets/README.md · ratings: calibrated by coordinator 2026-10-04 (from H17 agent (suggested))

RESULTS VS PREDICTIONS

Pending.

GOAL PERIODS IN DETAIL

G10 × failed (CK fails; t2* not beyond the sojourn null)
 G11 × failed (CK fails; t2* not beyond the sojourn null;
 MSM does not beat R1/M0)
 G12 ~ mixed (CK fails)
 G13 × failed (CK fails; t2* not beyond the sojourn null)
 G16 ~ mixed (CK fails)
 G17 ~ mixed (CK fails)
 G18 ~ mixed (CK fails)
 G19 ~ mixed (CK fails)
 G20 × failed (CK fails; t2* not beyond the sojourn null)
 G21 ~ mixed (CK fails)
 G23 × failed (CK fails; t2* not beyond the sojourn null)
 G24 × failed (CK fails; t2* not beyond the sojourn null)
 G25 ~ mixed (CK fails)
 G26 × failed (CK fails; t2* not beyond the sojourn null)
 G27 ~ mixed (CK fails)
 G30 × failed (CK fails; t2* not beyond the sojourn null)
+11 more periods: see the card.

FAITHFULNESS SCORECARD

A mapping	1	F identifiability	1
B assumptions	0	G ground truth	1
C adequacy	1	H comparative	1
D unfitted predictions	0	I transfer	0
E interventional	0		

Level: **hypothesis**. Scale 0 not done or failed, 1 partial, 2 passed (`writeup/paper.tex`).

CAVEATS

Pending.

ROUND 2

- **Essence.** What are the metastable cognitive modes (planning, building, debugging, existential talk), and what switches them?
- **H17-R1** A soft Jev-state MSM recovers slow cognitive modes, and switches are driven by external input.
- **H17-R2** Debugging is the slowest mode and the main sink of time.
- **H17-R3** Mode switches synchronize across agents only through shared artifacts (E3).

H17-behavior-metastable-sets · generated 2026-10-03 19:00 by infra/summaries/build-summaries.py · card hypotheses/H17-behavior-metastable-sets/README.md · ratings: calibrated by coordinator 2026-10-04 (from H17 agent (suggested)) (page 2)