


H23 · Distilling the village into a fine-tuned leader copies its vocabulary but transforms its plans  
Exploratory round 1 (G44, non-holdout), 2026-10-03. #45 confirmation written, not run.


Pipeline (from the record): the deployed leader (kimi-leader-v7-aug-64; agent 28 in #44, agent 30 in #45) is base Kimi K2.6  
LoRA-fine-tuned on 64 rows: 51 base-Kimi self-distilled replies to 20 synthetic #best scenarios + 13 agent-written rows.  
No village chat was a training target. 31/64 rows recovered verbatim from the agents' terminals.

P1 vocabulary copied: marker rate 14.2 vs 10.0 /100 tok (p=0.011) SUPPORTED; bigram copy 0.064 vs 0.055 INCONCLUSIVE.  
P2a plans not copied: directive share 0.62 (hand 0.25) vs corpus 1.00 SUPPORTED.  
P2b plans transformed: excess plan copy -0.04 (perfect copy 0.23), I\_transform within null → NOT TRANSMITTED.  
P3 closer to corpus than base: d 0.094 vs controls 0.096 FAILED; Kimi field invariance FAILED.  
Rivals: base echo (R1) not rejected; latest-chat echo (R2) fits the plans better than the corpus (JSD 0.15 vs 0.35).


**Corpus-distinctive n-grams**




**Corpus bigram copy**




**Marker classes**




**Plan profiles (primary act)**




**Plan copy: observed vs perfect-copy (hatched)**



**Embedding: corpus vs base**



**Synthetic: plan channel power**  
(-- n=16, — n=70; ○ coder 0.6, △ perfect)



**Synthetic: marker test power**  
(dotted: observed ratio)

