


H27: Critical slowing down warns of herding waves ahead

Vivian Rogers, with Claude Opus 5.5 · agent-swarm physics · Updated 2026-10-04

Complete 50% Faithfulness 1/5 Usefulness 2/5


Bottom line. Critical slowing down does not warn of herding in the AI Village: pile-ons in attention and in commits are fast, usually marked by a link, and too few agents are watched for slowing down to show.

level: hypothesis · pipeline checks 86% · updated 2026-10-04



WHAT HAPPENED

A pre-registered rule on project labels defined herding onsets: a project's share of labeled agents jumps from ≤ 0.25 to ≥ 0.5 within an hour and stays there. It found 21 onsets in 14 non-holdout weeks at 15-min resolution, mostly in shared-artifact weeks. Early-warning trends (rolling lag-1 autocorrelation and SD over the previous 6 h) were compared with matched placebo segments; 8 onsets were evaluable. The composite separated them weakly at 15 min (AUC 0.62) and not at 30 min; autocorrelation never rose (0.41), so this is not critical slowing down. **Round 1b** (deterministic labels, work ledger) reproduces all of it: same onsets, same AUC. The frozen alarm now beats a time-shifted copy at $p = 0.038$ (was 0.07), but a naive share alarm matches its 5/21 hits with fewer alarms. Work-commit onsets coincide with attention onsets, and #40's pile-on ran in minutes after a link. Holdout not run.



Observables. (a) AUC of each early-warning trend, onset vs placebo segments, at 1 h lead. Series: a synthetic fold at village size (purple), a clean $N = 200$ fold (green), and the village at 15 min (blue, 8 onsets) and 30 min (orange, 6). Autocorrelation never discriminates; variance does only at 15 min, with wide CIs. (b) Onsets warned within 3 h against false alarms per active day. Open circles: the same alarms shifted in time. No alarm beats its shifted copy.

SO WHAT? PRACTICAL VALUE FOR UNDERSTANDING AGENT SWARMS

For an operator this is a useful negative, now checked on cleaner data and in commits. Generic early-warning indicators (rising autocorrelation and variance) do not forecast pile-ons in a 10–15-agent LLM swarm. Pile-ons are fast pushes: a link posted in chat, then half the room within minutes (#40: 4 min to half the room's attention, 19 min to half its commits). A 6-hour trend window cannot see that, and at $N \approx 12$ a true fold would be barely detectable anyway. So don't build a critical-slowness dashboard for swarms this size. Watch announcements and shares instead, and size any response to minutes, not hours. The transferable lesson: test alarm timing against a rate-matched shift; the round-1 vs 1b p (0.07 vs 0.04) shows how thin such timing signals are.

PHYSICAL MODEL

Models 10 (kinetic mean-field Potts) and 01 (mean-field susceptibility). σ_i : agent i 's project in a window ($0 =$ rare projects, uncoupled). $k_a^{(-i)}$: other agents on a ; $h_a(t)$: a slowly drifting popularity field; βJ : uniform herding coupling; γ : reconsideration rate; x_a : share on a . Near a fold of x_A the leading eigenvalue $\lambda \rightarrow 0$, so the lag-1 autocorrelation ρ_1 and the variance rise (critical slowing down). Rivals (field steps, noise-induced flips) give no warning.

$$P(\sigma_i \rightarrow a) \propto e^{h_a(t) + \beta J k_a^{(-i)} / (N-1)} \quad (a \geq 1), \quad \propto 1 \quad (a = 0)$$

$$\dot{x}_a = \gamma [\text{softmax}_a(h(t) + \beta J x) - x_a]$$


$$\rho_1 = e^{-\gamma |\lambda| \Delta t} \rightarrow 1, \quad \text{Var } x_A \propto 1/(N|\lambda|) \quad (\lambda \rightarrow 0)$$

$$\text{alarm: } \tau_K(\hat{\rho}_1) > \tau^* \wedge \tau_K(\hat{\sigma}) > \tau^*, \quad \tau^* = 0.54$$

H27-herding-early-warning · generated 2026-10-04 01:33 by infra/summaries/build_summaries.py · card hypotheses/H27-herding-early-warning/README.md · ratings: RE-P1 round-1b agent (suggested; previous calibrated ratings 40 / 1.0 / 2.0)

RESULTS VS PREDICTIONS

Prediction	Observed (round 1 → 1b)	Verdict
P0: 10–25 onsets	21 → 21; 8 evaluable	yes
P1: $AUC \geq 0.70$	0.62 → 0.62 [0.50, 0.76]; 30 min 0.44 → 0.50	mixed
P2: $SD > AR1$	0.70/0.42 → 0.71/0.41	partly
P3: no operator value	5/21 hits; shift p 0.07 → 0.038; mostly level alarm 5/21	
P4: synthetic	$N=200$: 0.75; $N=12$: 0.58	2 of 4
R1b-1 replication	same onsets, same verdicts	yes
R1b-2 work on-sets	8 vs 14; lag 0 (3) or +1 (1) window	mostly
R1b-3 links	11/21; 9 to known projects, 2 first links	yes
#40 minute clock	t_{50} 1.5 min after link; work 15 min later	2 of 3



Second figure. (a) AUC of each trend, onset vs placebo, at 1 h lead: round 1 (orange) vs round 1b shared labels (blue). Nothing moves; autocorrelation stays below 0.5. (b) #40 at the minute clock: share of the room with a first mention of the hub (orange) and a first commit to it (blue) after the kickoff. Half the room is in within one 15-min window (shaded).

GOAL PERIODS IN DETAIL

G17 d 1b: descriptive (unchanged)
 G18 d 1b: descriptive (unchanged)
 G19 d 1b: descriptive (unchanged)
 G20 d 1b: descriptive (unchanged)
 G24 d 1b: descriptive (unchanged)
 G25 d 1b: descriptive (unchanged)
 G26 d 1b: descriptive (unchanged)
 G30 ~ 1b: mixed (AUC 0.55, 2/4 hit; unchanged)
 G31 ~ 1b: mixed (native: work onsets coincide with attention)
 G33 d 1b: descriptive (unchanged)
 G35 d 1b: descriptive (unchanged)
 G37 d 1b: descriptive (unchanged)
 G38 d 1b: descriptive (unchanged)
 G39 d 1b: descriptive (unchanged)
 +5 more periods: see the card.

FAITHFULNESS SCORECARD

A mapping	1	F identifiability	1
B assumptions	0	G ground truth	1
C adequacy	0	H comparative	0
D unfitted predictions	0	I transfer	0
E interventional	0		

Level: **hypothesis**. Scale 0 not done or failed, 1 partial, 2 passed (`writeup/paper.tex`).

CAVEATS

21 onsets, 8 evaluable, 4 in work: per-week AUCs rest on 1–2 onsets. Many indicators, leads, window lengths and two onset rules; only the 15-min composite at 1 h was primary. The alarm-vs-shift p crosses 0.05 between null draws (0.038–0.07): treat it as borderline. The step rule misses slow rises; the slow-rise variant gives 0.74 at 15 min. The link association is post hoc; most linked onsets are bursts of known projects (H28). Work labels are public agent commits (dense from #30). #51 supplies most false alarms.

ROUND 2

- **Essence.** Can an operator see a pile-on coming early enough to act, and from what signal?
- **H27-R1** Model onsets as announcement-triggered cascades: hazard of a majority pile-on after a link post, by poster, room size and the project's prior share; joint with H28.
- **H27-R2** Use a finer clock: per-turn project states from actions (minutes, not 15-min windows) so 6 h of history is not needed and the first 1.5 days become evaluable.
- **H27-R3** Test the one regime where the theory says warnings can work: long, large swarms with slowly drifting incentives (#51 roles, if any shared project ever rises), with the slow-rise onset rule.