


H35: The nudge is a measurably inefficient Maxwell demon

Vivian Rogers, with Claude Opus 5.5 · agent-swarm physics · Updated 2026-10-03

not yet rated

Bottom line. *Pending.*

level: hypothesis · pipeline checks 57% · updated —



WHAT HAPPENED

Not yet summarized. The question: Bits of agent-state information used per extra active minute gained (H04/H16 effects), against an optimal-timing benchmark (nudge at the pause gate, once). This is the information-to-work efficiency of operator feedback, and gives a target for a better nudging policy. *Check:* mutual information between nudge timing and agent state vs the activity gained; policy counterfactuals on logged gates. Practical aim (usefulness-first batch): produce a number or rule an operator of an LLM agent swarm could compute from logs and act on.

PHYSICAL MODEL

From: 04 (Kolchinsky), 02. The round-1 agent specifies the exact variant.

SO WHAT? PRACTICAL VALUE FOR UNDERSTANDING AGENT SWARMS

Pending.

Observables. *Pending.*

RESULTS VS PREDICTIONS

Pending.

FAITHFULNESS SCORECARD

A mapping	–	F identifiability	–
B assumptions	–	G ground truth	–
C adequacy	–	H comparative	–
D unfitted predictions	–	I transfer	–
E interventional	–		

Level: **hypothesis**. Scale 0 not done or failed, 1 partial, 2 passed (`writeup/paper.tex`).

CAVEATS

Pending.

GOAL PERIODS IN DETAIL

No goal-period folders yet.

ROUND 2

No round-2 redirects yet.
H35-nudger-maxwell-demon · generated 2026-10-03 19:00 by in-
fra/summaries/build-summaries.py · card hypotheses/H35-nudger-maxwell-
demon/README.md · ratings: not rated · DRAFT: auto stub, not yet written
(page 2)