


H40: The call clock sets the coupling

Vivian Rogers, with Claude Opus 5.5 · agent-swarm physics · Updated 2026-10-04

Complete 45% Faithfulness 2.5/5 Usefulness 3/5


Bottom line. In the computer-use scaffold agents couple per model call, not per minute, so per-hour coupling is per-call coupling (a family trait) times call rate and slow-cadence agents are heavy spins; the regime-I chat scaffold breaks this.

level: hypothesis · pipeline checks 86% · updated 2026-10-04



WHAT HAPPENED

Do agents respond to each other per model call or per minute? On 33 non-holdout periods I fitted a reply-hazard model on each recipient's real call schedule (context ledger, DQ2 reply labels), after checking on simulated replies that it tells the two clocks apart. In the computer-use scaffold (regimes II–III, 11 periods) coupling runs on the call clock: a call that spans more wall time carries no more reply hazard ($\eta = 0.00 \pm 0.05$ in G51; -0.03 at timer wakes fixed in advance). In G51 per-call coupling is flat in cadence ($s = +0.11 [-0.03, 0.25]$) and set by model family, while per-hour coupling scales with call rate (slope 1.0): slow agents are heavy spins. In regime I (22 periods) it fails: $\eta = +0.51 [0.41, 0.61]$. Natives: G51 yes, NE41 mixed, G18 and NE14 no. Holdout: not run.



Observables. (a) η per goal period: regime-I periods sit between the clocks, regime II–III at the call clock (0). (b) G51, 243 agent×unit points: per-call coupling (blue, model intercepts) is flat in call rate; per-hour coupling (orange, intercept + $\ln r$) rises with slope ≈ 1 . Lines: fitted slopes.

SO WHAT? PRACTICAL VALUE FOR UNDERSTANDING AGENT SWARMS

Useful, with a regime caveat. In a computer-use agent loop, an agent's wall-time responsiveness to others is its per-call coupling (a model-family trait) times how often it calls the model. Long tool calls, long pauses and slow models make agents "heavy": in G51 the slowest fifth answer within 5 minutes $8\times$ less often than the fastest. Halving call intervals raises 5-min coupling by about 50% (G51 32%), less at 30 min. So cadence is a coupling knob an operator can turn without touching prompts, e.g. short pause timers for coordinators, or long batch jobs to decouple an agent. It does not transfer to a scheduler-driven chat loop (regime I), where time itself carries part of the coupling. Mostly observational; the clean cadence intervention (NE20) is in the holdout.

PHYSICAL MODEL

Models 02 (kinetic Ising, Glauber attempt rates) and 09. Each model call c of recipient i is an update attempt; after reading message m at call 1, i replies at call n with hazard h_n . α_i : per-call coupling. e_n : wall time call n spans; a_n : age of m ; b_n : newer messages (burial); k_n : batch size. η : exposure-time elasticity (0 call clock, 1 wall clock). r_i : call rate. $\varepsilon(T)$: cadence elasticity of replying within T .

$$\text{cloglog } h_n = \alpha_i + \phi \ln n + \psi \ln a_n + \chi \ln(1+b_n) + \eta \ln e_n + \delta \ln(1+k_n) + \dots$$

$$\text{wall clock: } h_n = 1 - e^{-\lambda e_n} \Rightarrow \eta = 1; \quad \text{call clock: } \eta = 0$$


$$J_i^{\text{hour}} = J_i^{\text{call}} r_i^{1-\eta}, \quad \alpha_i = c + s \ln r_i \quad (s = 0 : \text{per-call invariant})$$

$$\varepsilon(T) = \partial \ln P(\text{reply} \leq T) / \partial \ln(\text{speed-up})$$

H40-call-clock-coupling · generated 2026-10-04 09:22 by infra/summaries/build_summaries.py card hypotheses/H40-call-clock-coupling/README.md · ratings: H40 agent (suggested)

RESULTS VS PREDICTIONS

Prediction	Observed	Verdict
P1 pooled $\eta \in [-0.25, 0.25]$	+0.30 [0.17, 0.43]; II -0.01, III -0.16, I +0.51	fail; II-III yes
P2 decay per call (III)	$\phi = -0.60, \psi = -0.19$; 6/8	yes
P3 $s \approx 0$, excl. -1	pooled -0.12 [-0.28, 0.04]; G51 +0.11	yes
P3 per-hour slope ≥ 0.5	III +0.74; I 0.00	III only
P4 held-out call > wall	33/33	yes
P5 collapse in calls	21/33	no (marginal)
P6 $\varepsilon(5) \in [0.5, 1]$, $\varepsilon(30)$ lower	0.57 [0.36, 0.78]; 0.47	yes
N1 G51 pause dose, heavy spins	$\eta_1 = -0.03$; R^2 cadence 0.50 vs family 0.33	yes
N2 NE41 no catch-up	read-out $\times 0.44$; ratio 1.86 (band fail)	mixed
N3 G18 / N4 NE14	$\eta = 0.81$ / 1 of 11 agents	no / no



Second figure. Synthetic replies on the real schedules of G18, G31, G38 and G51c, with call starts jittered within their bounds. $\hat{\eta}$ recovers 0, 1 and 0.5 (bias ≤ 0.05 , power 1.0 against the wrong clock), and is unmoved by agent couplings correlated with cadence (S4) or detection loss (S5). The between-agent slope is recovered (S4: 0.5) but noisy per period.

GOAL PERIODS IN DETAIL

G02 $\sim$ mixed
G03 $\times$ failed
G04 $\sim$ mixed
G05 $\sim$ mixed
G06 $\sim$ mixed
G08 $\sim$ mixed
G11 $\sim$ mixed
G12 $\times$ failed
G13 $\checkmark$ supported
G16 $\sim$ mixed
+24 more periods: see the card.

FAITHFULNESS SCORECARD

A mapping	1	F identifiability	2
B assumptions	1	G ground truth	1
C adequacy	1	H comparative	1
D unfitted predictions	1	I transfer	1
E interventional	1		

Level: **hypothesis**. Scale 0 not done or failed, 1 partial, 2 passed (`writeup/paper.tex`).

CAVEATS

Regime I is unexplained: positive η survives logged start times and short spans, but endogenous generation time is untested. Cadence is partly chosen: only the pause-timer dose is quasi-exogenous. The cross-sectional G51 slope (1.0) exceeds the model's speed-up elasticity (0.40): composition, not only cadence. Reply labels are DQ2 primary parents (one per message), which undercounts multi-target replies, most for big batches. The NE41 ratio failed for recency reasons found post hoc. Two post hoc checks; no holdout run.

ROUND 2

Run `confirm.py` (NE20 one-tool-call DiD is the cleanest cadence intervention). Decompose regime-I chat spans into generation and scheduler wait. Repeat on the content channel (H29 boundary pull on the ledger).

- **Essence.** a measured coupling-per-call that an operator can multiply by cadence, and an explanation of why the regime-I chat scaffold does not obey it.
- **H40-R1** Why regime I is not a call clock.
- **H40-R2** Content channel.
- **H40-R3** Lags in call cycles (HH187).
- **H40-R4** Run `confirm.py` on NE20/NE44 and the transfer periods; NE20 is the cleanest cadence intervention in the dataset.
- **H40-R5** Fix the collapse test by matching agents on per-call coupling (family) before comparing cadence tertiles, and add multi-parent reply labels (DQ2 holds only the primary parent).

H40-call-clock-coupling · generated 2026-10-04 09:22 by infra/summaries/build-summaries.py · card hypotheses/H40-call-clock-coupling/README.md · ratings: H40 agent (suggested) (page 2)