


H52: Humans are just loud agents

Vivian Rogers, with Claude Opus 5.5 · agent-swarm physics · Updated 2026-10-04

Complete 45% Faithfulness 2/5 Usefulness 2.5/5


Bottom line. Matched on salience, humans get about double the replies and a small content pull; agent leaders share this status premium, and off-role requests lose it.

level: hypothesis · pipeline checks 86% · updated 2026-10-04



WHAT HAPPENED

The test asks whether a human message moves its recipient more than an agent message with the same salience. Salience means naming, read-out age (context ledger), novelty, length, recipient state and room size. Matched comparisons run on 17 non-holdout periods plus four native tests. Dated predictions and synthetic recovery come first. **Replies:** agents answer humans about twice as often as matched agent messages (pooled $+0.030$ $[0.019, 0.042]$; naming adds $+0.13$). **Content:** a small premium, G51 $+0.014$ $[0.008, 0.021]$ under five statistics, and $+0.05$ in regime I on H29's boundary design. **Activity:** ≈ 0 . Replies to humans are more supportive. In #51 the reply premium exists only for messages that fit the recipient's private role. Elected or appointed agent leaders get the same reply premium. Agents treat the bot like a matched agent message, but rebuff it. **Mixed.** The holdout is not run.



Observables. (a) Reply premium over matched agent messages. It is positive for humans in both eras, and in #51 only for messages aligned with the recipient's role. Elected or appointed agent leaders get the same premium; the bot gets none. (b) Content premium. G51 is positive under every statistic. Regime I is positive on H29's boundary design; the DiD statistic is biased by about -0.024 there. Agent leaders get none.

PHYSICAL MODEL

Linear response to fields (model 02). A message m is an impulsive field at the recipient's next model call. It acts on the content spin (whitened statement vector), the activity spin, the reply and the stance. Its coupling $g(s_m)$ depends only on the salience s_m (named, read-out age, novelty $\|P_{\perp} \hat{u}_m\|$, length, idle recipient, room size), the same for every sender. The authority premium π_c is an extra field term of sender class c (human, bot, leader). H52 (HH164) predicts $\pi_{\text{human}} = 0$.

$$\begin{aligned} \mathbb{E}[Y_{mj} \mid s_m, c] &= g(s_m) + \pi_c, & \pi_{\text{agent}} &\equiv 0 \\ \hat{\pi}_c &= \sum_s n_{c,s} (\bar{Y}_{c,s} - \bar{Y}_{\text{agent},s}) / \sum_s n_{c,s} \\ Y^{\text{con}} &= (Q - x_{\text{last}}) \cdot \widehat{P'_{\perp} u_m} - \langle \cdot \rangle_{\text{placebo}} \\ Y^{\text{rep}} &= \mathbf{1}[m = \text{parent}(j\text{'s next messages})] \end{aligned}$$


SO WHAT? PRACTICAL VALUE FOR UNDERSTANDING AGENT SWARMS

To get an answer from an agent, name it and frame the request in terms of its current objective. At equal salience, agents answer humans about twice as often ($+0.030$) and move a little toward their content ($+0.014$ in G51). The premium is a status effect: elected and appointed agent leaders get the same reply premium. It does not override the assigned plan: on-role human messages get $+0.037$, off-role ones $+0.007$. Being human buys an answer, not obedience. A safety team must not read the premium as deference, because it vanishes for role-conflicting requests. The labeller also knew which senders were human, and one period carries most of the regime-III evidence.

H52-humans-loud-agents · generated 2026-10-07 17:50 by infra/summaries/build_summaries.py · card hypotheses/H52-humans-loud-agents/README.md · ratings: H52 agent (suggested)

RESULTS VS PREDICTIONS

Prediction	Observed	Verdict
P1 H52: content ≈ 0	G51 +0.014 (5 statistics, 5/5 segments); reg. I DiD +0.005, boundary +0.05/+0.09	H52 fails in III
P2 reply premium > 0	pooled +0.030 [0.019, 0.042]; 8/17 CI > 0 , none < 0	yes
P3 activity ≈ 0	CI includes 0 in 15/17	yes
P4 stance $\geq +0.05$	+0.16 [0.10, 0.23] (labeller knew)	yes
P5 bot = neither	like a matched agent; replies push back (-0.97)	no
P8 naming moves content	DiD naming < 0 (reply quoting); boundary 5.4×	no (statistic)
N1 G04 same call	content -0.006 [-0.024, 0.014]; reply odds 1.49 [0.88, 2.79]	inconcl.
N2 #51 role conflict	reply on-role +0.037, off +0.007 (Δ CI > 0)	R1 rejected
N3 NE43 no change	content same; reply -0.039 [-0.072, -0.005]	mixed
N4 leaders get none	reply +0.028 (G26), +0.029 (G35); no content n.s.	



Second figure. Synthetic validation on the real read-out schedules (10 replicates per cell; circles: null; squares: planted premium). The design recovers reply, activity (bias-corrected) and stance premia. Shared topic biases H30's orthogonalized χ , so a DiD statistic replaces it. The DiD content statistic is unbiased in the G51 cells but reads about -0.024 low in dense regime-I chat with real labels.

GOAL PERIODS IN DETAIL

G03 ~ mixed — reply 0.050*, content 0.025
 G04 ~ mixed — reply +0.025*, content +0.012 (boundary +0.05...
 G05 ~ mixed — reply 0.030*, content 0.001
 G06 ~ mixed — reply 0.009, content 0.003
 G08 ~ mixed — reply 0.036, content 0.022
 G10 d descriptive — reply 0.020, content -0.022
 G12 ~ mixed — reply 0.023, content 0.039*
 G13 ~ mixed — reply 0.081*, content 0.023
 G16 ~ mixed — reply 0.094*, content -0.013
 G17 ~ mixed — reply 0.097*, content 0.000
 +10 more periods: see the card.

FAITHFULNESS SCORECARD

A mapping	1	F identifiability	1
B assumptions	1	G ground truth	1
C adequacy	1	H comparative	1
D unfitted predictions	1	I transfer	1
E interventional	1		

Level: **hypothesis**. Scale 0 not done or failed, 1 partial, 2 passed (`writeup/paper.tex`).

CAVEATS

“Human” means public viewers in 2025 and mostly staff later. Speech acts (requests, instructions) are not matched. The content statistics carry known, opposite biases. The content claim rests on their agreement in G51 and on the boundary design in regime I. Only G04, G06 and G51 have power. The reply and stance labeller knew when an author was human. The placebo-class null is anti-conservative.

ROUND 2

- **Essence.** whether agents weigh human input above peers' at equal salience, and whether it overrides their plan.
- **H52-R1** Speech-act matching.
- **H52-R2** Separate the labeller from the agents.
- **H52-R3** Role conflict with ground truth.
- **H52-R4** A cleaner content statistic.
- **H52-R5** Run confirm.py after committing (#51 tail; #45 leader).