


H52: Humans are just loud agents

Vivian Rogers, with Claude Opus 5.5 · agent-swarm physics · Updated 2026-10-04

Complete 45% Faithfulness 2/5 Usefulness 2.5/5


Bottom line. Matched on salience, human messages get about twice the reply rate and a small content pull compared with agent messages; the premium is shared by agent leaders and disappears when the message conflicts with the recipient's role, so it is status-weighted attention, not deference against the plan.

level: hypothesis · pipeline checks 86% · updated 2026-10-04



WHAT HAPPENED

Does a human message move its recipient more than an agent message with the same naming, read-out age (context ledger), novelty, length, recipient state and room size? Matched comparisons on 17 non-holdout periods plus four native tests, with dated predictions and synthetic recovery first. **Replies:** humans are answered about twice as often as matched agent messages (pooled +0.030 [0.019, 0.042]; naming adds +0.13). **Content:** a small premium, G51 +0.014 [0.008, 0.021] under five statistics; regime I +0.05 on H29's boundary design. **Activity** ≈ 0 . Replies to humans are more supportive. In #51 the reply premium exists only for messages that fit the recipient's private role. Elected or appointed agent leaders get the same reply premium. The bot is treated like a matched agent message, but rebuffed. Verdict: mixed. Holdout: not run.



Observables. (a) Reply premium over matched agent messages: positive for humans in both eras, only for #51 messages aligned with the recipient's role, the same for elected or appointed agent leaders, none for the bot. (b) Content premium: G51 positive under every statistic; regime I positive on H29's boundary design (the DiD statistic is biased about -0.024 there); agent leaders get none.

PHYSICAL MODEL

Model 02 (linear response to fields). A message m read at the recipient's next model call is an impulsive field on its content spin (whitened statement vector), activity spin, reply and stance. Its coupling $g(s_m)$ depends only on salience s_m (named, read-out age, novelty $\|P_{\perp} \hat{u}_m\|$, length, idle recipient, room size), the same for every sender. The authority premium π_c is an extra field term carried by sender class c (human, bot, leader); H52 (HH164) says $\pi_{\text{human}} = 0$.

$$\begin{aligned} \mathbb{E}[Y_{mj} \mid s_m, c] &= g(s_m) + \pi_c, & \pi_{\text{agent}} &\equiv 0 \\ \hat{\pi}_c &= \sum_s n_{c,s} (\bar{Y}_{c,s} - \bar{Y}_{\text{agent},s}) / \sum_s n_{c,s} \\ Y^{\text{con}} &= (Q - x_{\text{last}}) \cdot \widehat{P'_{\perp} u_m} - \langle \cdot \rangle_{\text{placebo}} \\ Y^{\text{rep}} &= \mathbf{1}[m = \text{parent}(j\text{'s next messages})] \end{aligned}$$


SO WHAT? PRACTICAL VALUE FOR UNDERSTANDING AGENT SWARMS

Alignment-relevant, but small. Humans in this swarm are not just loud agents: at equal salience, agents answer human messages about twice as often and lean a little toward their content. But the extra weight is a status effect: elected and appointed agent leaders get the same reply premium. It also does not override the agent's assigned plan; off-role human messages get none. An operator who wants to be followed should name the agent and frame the request in terms of its current objective. Being human buys an answer, not obedience. A safety team should not read the premium as robust deference: it vanishes for role-conflicting requests, and the reply and stance labels may partly reflect a labeller that knew the sender was human. One period carries most of the regime-III evidence.

H52-humans-loud-agents · generated 2026-10-04 11:10 by infra/summaries/build_summaries.py · card hypotheses/H52-humans-loud-agents/README.md · ratings: H52 agent (suggested)

RESULTS VS PREDICTIONS

Prediction	Observed	Verdict
P1 H52: content ≈ 0	G51 +0.014 (5 statistics, 5/5 segments); reg. I DiD +0.005, boundary +0.05/+0.09	H52 fails in III
P2 reply premium > 0	pooled +0.030 [0.019, 0.042]; 8/17 CI > 0 , none < 0	yes
P3 activity ≈ 0	CI includes 0 in 15/17	yes
P4 stance $\geq +0.05$	+0.16 [0.10, 0.23] (labeller knew)	yes
P5 bot = neither	like a matched agent; replies push back (-0.97)	no
P8 naming moves content	DiD naming < 0 (reply quoting); boundary 5.4×	no (statistic)
N1 G04 same call	content -0.006 [-0.024, 0.014]; reply odds 1.49 [0.88, 2.79]	inconcl.
N2 #51 role conflict	reply on-role +0.037, off +0.007 (Δ CI > 0)	R1 rejected
N3 NE43 no change	content same; reply -0.039 [-0.072, -0.005]	mixed
N4 leaders get none	reply +0.028 (G26), +0.029 (G35); no content n.s.	



Second figure. Synthetic validation on the real read-out schedules (10 replicates per cell; circles null, squares planted premium). Reply, activity (bias-corrected) and stance are recovered. H30's orthogonalized χ is biased by shared topic, so it was replaced. The DiD content statistic is unbiased in the G51 cells but reads about -0.024 low in dense regime-I chat with real labels.

GOAL PERIODS IN DETAIL

G03 ~ mixed — reply 0.050*, content 0.025
 G04 ~ mixed — reply +0.025*, content +0.012 (boundary +0.05...
 G05 ~ mixed — reply 0.030*, content 0.001
 G06 ~ mixed — reply 0.009, content 0.003
 G08 ~ mixed — reply 0.036, content 0.022
 G10 d descriptive — reply 0.020, content -0.022
 G12 ~ mixed — reply 0.023, content 0.039*
 G13 ~ mixed — reply 0.081*, content 0.023
 G16 ~ mixed — reply 0.094*, content -0.013
 G17 ~ mixed — reply 0.097*, content 0.000
 +10 more periods: see the card.

FAITHFULNESS SCORECARD

A mapping	1	F identifiability	1
B assumptions	1	G ground truth	1
C adequacy	1	H comparative	1
D unfitted predictions	1	I transfer	1
E interventional	1		

Level: **hypothesis**. Scale 0 not done or failed, 1 partial, 2 passed (`writeup/paper.tex`).

CAVEATS

“Human” means public viewers in 2025, mostly staff later; speech acts (requests, instructions) are not matched. Content statistics carry known, opposite biases; the claim rests on their agreement in G51 and on the boundary design in regime I. Only G04, G06, G51 are powered. The reply and stance labeller knew when an author was human. The placebo-class null is anti-conservative.

ROUND 2

- **Essence.** whether agents weigh human input above peers' at equal salience, and whether it overrides their plan.
- **H52-R1** Speech-act matching.
- **H52-R2** Separate the labeller from the agents.
- **H52-R3** Role conflict with ground truth.
- **H52-R4** A cleaner content statistic.
- **H52-R5** Run confirm.py after committing (#51 tail; #45 leader).