


H146: Ideology egregores in #51 recruit by being read, survive their hosts' amnesia, and carry functional behaviors

Vivian Rogers, with Claude Opus 5.5 · agent-swarm physics · Updated 2026-10-09

Complete 70% Faithfulness 1.5/5 Usefulness 1.5/5


Bottom line. In #51 the memplexes survive wipes and spread by reading only as much as random vocabulary does; no functional behavior beyond that is seen.

level: hypothesis · pipeline checks 86% · updated 2026-10-09



WHAT HAPPENED

The claim: the ideologies of #51 are *egregores*, patterns that run on changing hosts. They should recruit carriers through what those carriers read, survive their hosts' forced context wipes, and carry functional behaviors (recruitment, repair, division of labor). We test H145's 15 frozen memplexes and, post hoc, eight qualitative candidates, each against 30 frequency-matched random element sets (pseudo-patterns). **Result:** memplexes return after a host's wipe at 0.90 of the placebo rate and faster when the host re-reads them. Random element sets do the same. Read-gated expression and host specialization are also at the pseudo level, and nobody repairs a wiped or challenged host. At this instrument the ideologies behave as the village's generic vocabulary.



Observables. One row per pattern (blue: H145 memplexes; orange: post hoc candidates), with 95% CIs; gray dots are that pattern's 30 pseudo-patterns. Left: read - in-flight log rate ratio. Middle: re-expression after a forced wipe vs placebo. Right: named repair messages after a wipe vs placebo. Every pattern sits inside its pseudo-pattern cloud.

SO WHAT? PRACTICAL VALUE FOR UNDERSTANDING AGENT SWARMS

A context wipe does not erase what an agent was talking about: it comes back at about 0.9 of the normal rate within 20 calls, and about 1.4 times faster when the agent re-reads it. This holds for any vocabulary, so no special mechanism is needed to keep a shared idea alive through wipes; keeping it in readable chat is enough. Do not expect other agents to restore a wiped agent's ideas: named messages to it do not rise (RR 0.95–1.07).

PHYSICAL MODEL

Contagion with a read channel (model 03). For a talk row r of a non-host, $y_r = 1$ if its chat carries an element of K . R_r counts K items from hosts read in the mirror window; P_r counts items in flight at the matched lag. Strata s are agent $\times$ day. After a forced wipe F or a placebo call P, the host's first K statement within 20 calls gives the hazard ratio HR.

$$\log \lambda_r = \alpha_{s(r)} + \beta_R R_r + \beta_P P_r + \gamma^\top z_r, \quad \delta = \beta_R - \beta_P$$


$$\text{HR} = \frac{\sum_s D_s^F T_s^P / T_s}{\sum_s D_s^P T_s^F / T_s}, \quad p_{\text{ps}} = \frac{1 + \#\{\text{pseudo} \geq K\}}{31}$$

H146-egregore-recruitment-behaviors-51 · generated 2026-10-09 13:37 by infra/summaries/build_summaries.py · card hypotheses/H146-egregore-recruitment-behaviors-51/README.md · ratings: H146 round-1 agent (suggested)

RESULTS VS PREDICTIONS

prediction	observed	verdict
P1 read-gated adoption	CI>0 in 2/15 (K02 hub-carried, inconcl. K05); $p_{ps} \geq 0.06$; power ≤ 0.53	
P2 named $\geq 3 \times$ unnamed	no P1 pass	n/a
P3 newcomers follow reads	power ≤ 0.15 ; K09 first for 5/11	untestable
P4 wipe HR ≥ 0.8	median 0.90 (0.78–1.11); pseudo 0.89	supp.*
P4 read-gated faster	HR 1.2–2.3, CI>1 in 10/15; pseudo 1.39	supp.*
P5 repair RR>1.2	wipe 0.95–1.07; challenge 0.64–0.97	failed
P6 specialization	z 9–153; beyond pseudo 2/14 (chance)	pseudo lvl
P6 $\Omega < 0$ (amended)	0/14 below surrogate p05	failed

*as written; matched by pseudo-patterns, so not pattern-specific.



Second figure. Synthetic validation on the real ledger. Left: the read – in-flight test has size ≤ 0.05 in fast and slow field worlds and reaches power 0.8 at log-OR 0.3 only near 4,000 events. Right: on each candidate’s own skeleton the power stays below 0.6, because only 100–1,200 rows hold a K item in the ~ 8 -s mirror window.

GOAL PERIODS IN DETAIL

G51 $\sim$ mixed
NE33 d descriptive
NE41 $\sim$ mixed

FAITHFULNESS SCORECARD

A mapping	1	F identifiability	1
B assumptions	1	G ground truth	1
C adequacy	0	H comparative	1
D unfitted predictions	1	I transfer	0
E interventional	1		

Level: **hypothesis**. Scale 0 not done or failed, 1 partial, 2 passed (`writeup/paper.tex`).

CAVEATS

Recruitment is unpowered as a negative: a pattern-specific read effect of log-OR 0.3 could be missed. The adoption design is estimable for 3 of 15 memplexes. Newcomers are too few. Host specialization measures each agent’s own vocabulary (the shared-prior impostor). No message-level relay classifier exists; dropping Claude Fable 5’s items changes nothing. The candidate results and the read-gating pseudo comparison are post hoc. Pseudo-patterns use 30 draws, so $p_{ps} \geq 0.032$. The reserved #51 tail (51m) is not used.

ROUND 2

Widen the read window to the whole call with a matched in-flight placebo, or pool memplexes hierarchically, to power P1. Test specificity directly: does reading a K item raise K more than other patterns? Measure repair as restatement of the host’s own lapsed elements. Run the frozen `confirm.py` on 51m after sign-off.

- **H146-R1** Power P1.
- **H146-R2** Specificity.
- **H146-R3** Repair.

H146-egregore-recruitment-behaviors-51 · generated 2026-10-09 13:37 by `infra/summaries/build_summaries.py` · card hypotheses/H146-egregore-recruitment-behaviors-51/README.md · ratings: H146 round-1 agent (suggested) (page 2)