

Agent-swarm physics: hypothesis compendium

Vivian Rogers, with Claude Opus 5.5

(Living status document, regenerated 2026-10-04 16:21. Not the paper.)

Two RevTeX pages per hypothesis follow; click a row to jump to it. Ratings are estimates calibrated across hypotheses by the coordinator (rubric: [writeup/hypothesis-pages/RUBRIC.md](#)). **Completion:** estimated progress of the research direction (idea 5%, round 1 done 35–50%, holdout run 55–65%, robustness and causal designs 70–85%, settled 90–100%).  **Credence** (scoring v2, [writeup/scoring/scoring-v2.pdf](#)): probability that the scored claim (gray line under each title) survives an independent confirmatory test at its stated scope, built from an evidence ladder; “!” marks a fragile claim (its headline changed between data versions). Red = low, gray = 0.5, blue = high. **EU** (estimated usefulness) = credence × value if true (0–5: whose decision changes, by how much, how generally); gray ≈ 1.5. Most claims carry two independent scores (coordinator plus a blind rater, log-odds mean; splits beyond $\Delta p > 0.25$ or $\Delta V > 1$ adjudicated). How credence and faithfulness are built: [writeup/credence-faithfulness.pdf](#). Nothing here is holdout-confirmed unless the page says so. The writeup figures (32 hypotheses, six themes) follow the summaries as an appendix. Credence 

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H01	Emergent superagents exist No coordination-defined group is an effective (KW) superagent; the persistent unit is agent + own artifact	45%	0.62	1.4	11
H02	Inferred couplings reflect real influence Activity-timing pairwise couplings sit at the null floor; regime-III ‘collective coupling’ was day edges	65%	0.94!	2.4	13
H03	Swarm activity is self-exciting, and its criticality is set by the go... Activity is subcritical everywhere ($n \approx 0.2$ – 0.4)	50%	0.86	1.7	15
H04	External forcing reshapes the response kernel, reversibly A first nudge adds ~1.2 active minutes per 30 (A30 1.16 [0.75, 1.57]) to its named target only, starting at the recei...	65%	0.49!	0.9	17
H05	Cutting the cross-room channel lowers entropy production, and rooms b... Rooms are coupled blocks for talk, and the coupling runs through reading: co-location adds nothing once ledger read c...	55%	0.88	2.2	19
H06	Free weeks show neutral cooperative dynamics with a cooperator core Agents state private sub-tasks everywhere; work concentrates on a shared repo when one exists (neutral copying at bes...	50%	0.80	1.4	21
H07	The RPG forks diverge from a common ancestor at a measurable rate After the 03-16 split the two RPG forks diverge gradually (69% / 81% of ancestor files identical at the end of #35, n...	50%	0.52	0.8	23
H08	Context is the coupling Responses are gated at the recipient’s next model call (read-out)	50%	0.91	2.9	25
H09	Agent swarms have an effective thermodynamics Agents leave idle at a timer gate; unaddressed chatter holds them there (OR 0.50) and a mention releases them (OR 3.1...	40%	0.56!	1.3	27
H10	Goals are Legendre pushes: a swarm’s response to an assigned goal is ... Goal kickoffs push content along the goal direction, but pre-kickoff fluctuations do not predict who moves (no linear...	50%	0.79	1.4	29
H11	Division of labor vs herding is the sign of a Potts coupling, set by ... Agents herd onto shared projects in attention and work; ownership disperses	50%	0.75	2.1	31
H12	Groupthink is dimensional collapse: few collective modes, shrinking p... Talk and content collective modes are real; the activity market mode was an artifact; low-PR weeks are restatement lo...	52%	0.77!	1.5	33

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H13	Model families carry their own fields, and couple by family Model-family content fields are writing style; behavior carries only a weak family trace (tooling-like)	50%	0.87	2.2	35
H14	Entropy production of behavior-state sequences: per-family arrows of ... Fine-action irreversibility is real for every agent and survives scaffold removal; the regime contrast was scaffold	50%	0.82!	0.6	37
H15	Semantic information through natural scrambles: which information kee... A forced wipe cuts work commits ~39% for ~10 turns ($\approx 10\%$ of a segment's output); saving more memory doesn't soften it...	55%	0.78!	2.7	39
H16	Agents get stuck in metastable traps and escape by Kramers kinetics; ... Within agent, escape from inactive spells and timer-pause chains ages instead of being memoryless (TS1r deep beta abo...	48%	0.70	2.1	41
H17	Behavior is a Markov state model with a few metastable sets; its mixi... Behavior is non-Markov; the slow split is acting vs waiting; no metastable stuck states	45%	0.61	0.6	43
H18	Attention dilutes as $1/k$: response to a message falls with the number... In #51, attention per sender dilutes with batch size when batch size is set exogenously by a timer wake ($k^{-0.50}$ [0.4...	48%	0.64	1.3	45
H19	One curve for all periods: per-period loop gains collapse onto a func... Per-period loop gains do not collapse onto one positive curve of attended partners per attention (x_{att}): the pre-reg...	50%	0.71	0.7	47
H20	Aging: day-to-day content autocorrelation depends on time since kicko... At day resolution swarm content does not age as strongly as $\mu = 0.5$: in #51 $\mu = 0.5$ aging is rejected (power $\zeta = 0.9...$	50%	0.72	0.7	49
H21	The debate week (#12) is a two-sublattice antiferromagnet In #12's ten debates, drafted opponents disagree on the merits in 45% of their replies (true rate 0.52–0.76) and team...	60%	0.70	1.1	51
H22	Private, conflicting goals make #51 a spin glass In #51 (51b–51d), conflicting private goals produce no antagonism in content or stance: reading a rival's or opposed...	50%	0.60	0.6	53
H23	Distilling the village into a fine-tuned leader copies its vocabulary... In #44 the deployed self-distilled leader's plan-act profile departs sharply from its 100%-directive training corpus ...	40%	0.37	0.4	55
H24	In the forecast week (#21), switching from independent drafts to comp... In the forecast week, reading a teammate's document moves an agent's statements toward that teammate; there is no agg...	50%	0.42!	0.8	57

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H25	A distance-to-criticality dial, T/T _c , per window and per channel Daily equal-time gains are subcritical (~ 0.1 – 0.16); content 'near-critical' is a shared field	50%	0.85	1.7	59
H26	The swarm is near-critical in what it says but subcritical in when it... Content and talk co-move within rooms at a loop gain of ~ 0.5 (not near-critical); activity has no room gain once the ...	50%	0.84!	1.5	61
H27	Critical slowing down warns of herding waves hours ahead Critical slowing down does not warn of herding onsets	50%	0.43	0.8	63
H28	Links are the contagion vector of herding Links co-move with project switches in attention and work ($HR \approx 2.6$) but mark bursts rather than cause them: recipien...	52%	0.81	2.0	65
H29	Driver nodes: where an operator message moves the whole swarm Content influence is address-gated (named 0.05–0.08 per message vs unnamed ≈ 0 in #51); content-pull driver rankings ...	50%	0.65	2.1	67
H30	An operator-susceptibility gauge χ_{op} : how steerable is the swarm tod... In regime III a nudge adds about one extra active minute for its leading-@ target from the receiving call (G51 0.98 [...)	50%	0.39!	1.0	69
H31	Consensus time scales with the interaction graph's spectral gap Consensus time doesn't follow $1/\lambda_2$ (constant ~ 4 h forecast); deadline votes settle at read-out speed	50%	0.55	1.4	71
H32	A net information current identifies de facto leaders Exposure-gated content transfer is real (17/32 periods; read beats unread at 60 s in 17/17), but about a third of it ...	50%	0.78	1.8	73
H33	The diversity–productivity curve is an inverted U: the swarm's operat... No diversity–productivity optimum on committed work; the #51 tail is model-dependent	50%	0.70	1.1	75
H34	Idea cascades follow a power law whose exponent reads off the distanc... Idea cascades are subcritical (R 0.09–0.40) and exposure-locked in 32/32 periods, but about half of the lock is conv...	50%	0.84	1.9	77
H35	The nudger is a measurably inefficient Maxwell demon Nudges buy glances (+0.08 P(active within 30 min) per first nudge), not committed work; switching the nudger on or of...	55%	0.61!	1.1	79
H36	A reorganization alarm: susceptibility and multi-information peak at ... Plain topic shift detects goal changes (AUC ~ 0.95) far better than the physics alarm (AUC 0.66–0.68, not better than ...	50%	0.80!	2.2	81

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H37	Conflict lives in stance, not topic: stance spins are antiferromagnetic A validated zero-shot disagreement flag separates assigned debate opponents from teammates in #12 (33/74 vs 0/60; str...	55%	0.65	1.0	83
H38	Joint silences are platform stalls Regime-III co-activation is 70–80% the runner’s start/stop schedule	50%	0.90!	2.9	85
H39	Catalysts vs. fields No lever is a pure catalyst: nudges both unstick (escape $\times 1.3$ –1.6) and tilt (idle/wait share -5 to -8 points); mentio...	50%	0.67	1.7	87
H40	The call clock sets the coupling In the computer-use scaffold (regimes II–III), reply coupling runs on the call clock ($\eta \approx 0$): per-hour coupling = per...	45%	0.72	2.5	89
H41	Read-out gating gives the swarm a light cone Inside rooms, novel items follow the logged light cone (acausal $\leq 4\%$, median 0.6%) and rooms cage spread (75–100% of ...	48%	0.70!	1.9	91
H42	Cross-excitation is a delayed step at the next read-out Once each agent’s own call clock and call class are modelled, talk cross-excitation from read messages is near zero (...)	42%	0.61	1.8	93
H43	Kicks leave a refractory window No refractory window beyond one read-out; only same-call kicks are wasted	40%	0.69	1.9	95
H44	Erasure makes agents busy but unproductive Forced erasure triggers orderly re-reading, costs 4–11% of output, breaks loops	42%	0.88	3.2	97
H45	Context homeostasis: each agent keeps its context near a set point Room share of context is passive; the scaffold bounds context	40%	0.82	1.6	99
H46	Style is a conserved charge Style drifts with context fill and resets at erasure, yet identifies agents better than content	42%	0.75	2.4	101
H47	Rooms set the coherence length Rooms bound content coherence when they do different things; inside rooms coherence follows conversation	40%	0.66	1.6	103
H48	The settling time is a mixing time on the read-out graph In broadcast rooms read-out coverage is far faster than post-kickoff content settling: 90% of room-mate pairs read ea...	40%	0.51	0.8	105

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H49	Edge-trimmed regime III is a dilute ferromagnet Edge-trimmed regime-III excess is a dense/shared-field mode, not a dilute ferromagnet	40%	0.81	1.2	107
H50	Field or coupling: transfer functions and lag tell which Activity co-movement is the scheduler's field; talk is a coupling gated at hop 1	45%	0.85	3.0	109
H51	One dial: an effective coupling collapses the phase diagram The read-out loop gain g_lag does not collapse the phase diagram: it explains 0 of 4 unfitted observables across 33 p...	42%	0.61	0.9	111
H52	Humans are just loud agents Humans get a status-weighted attention premium (~2× replies), not deference	45%	0.60	2.1	113
H53	Herdling is announcement-seeded nucleation Herdling is seeded by a project's first chat link; deadline votes are read-out-bound	42%	0.59	1.6	115
H54	The kickoff text is the quench target Day-1 content lands on what the kickoff names; agents land on private goals	45%	0.92	2.6	117
H55	Norm-enforcers are the swarm's immune cells A read correction does not end blocked spells ($\Delta +0.007$ [-0.053, +0.070], 337 treated windows, power 0.96), and agent...	40%	0.57!	0.6	119
H56	Entropy production fingerprints the platform Turn-level EP does not detect platform changes; log features do	40%	0.69	1.6	121
H57	Copying beats transformation when the backlog is large A large read backlog does not make agents copy what they read: the raw read-set echo, an upper bound, rises by at mos...	40%	0.57	1.7	123
H58	Effective superagents are coordinated agents plus their artifacts In the powered units (#51b-d) no coordination-defined composite of agents plus artifacts predicts members' next artif...	40%	0.71	1.8	125
H59	One lever model for all operator inputs No single (field, catalysis, read-out delay) triple transfers across operator input classes (leave-one-class-out fail...	40%	0.59	1.2	127
H60	An index policy beats the nudger No index policy beats a simple rule: nudge each trap once at its second idle gate ($\times 1.55$ [1.25, 1.81] over the logged...	55%	0.44	0.7	129

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H61	Contagiousness is predictable at first use An idea's spread is mostly not forecastable at first use beyond who posted it (seed features beat the class + poster ...	45%	0.64	1.0	131
H62	Ideas travel the reply graph, not the room Reply-partner exposure is followed by 2–3× more idea adoption per edge than room-only exposure (A 1.92 [1.52, 2.42], ...	48%	0.70	1.4	133
H63	Herding bursts start with a work signal Herding bursts are announced, not caused by work: a chat link precedes the second agent's arrival in 60% of bursts vs...	40%	0.56	1.0	135
H64	Conflict needs a scarce prize Settling the prize switches assigned conflict off: in #12 the validated disagreement flag marks 44% of opponent repli...	45%	0.64	1.0	137
H65	Leaders are routers, not sources Designated leaders are attention hubs (+0.39 to +0.45 replies per statement), not content routers; attention is unali...	45%	0.63	1.3	139
H66	Shared platform latency is a common-noise field A shared platform latency field exists in #44 and #51 (11/29 units) and grows with N, but it is load, not drive: it r...	45%	0.58	0.9	141
H67	A lagged criticality dial The read-out loop gain of talk (extra messages one message causes through reading at hop 1) is subcritical in 66/66 u...	45%	0.79	2.2	143
H68	Attention dilution is a mixture of two strategies Attention dilution is one narrow unimodal exponent per period (between-agent SD 0.09–0.20), not a mixture of thin att...	40%	0.76	1.3	145
H69	Restatement loops are context fixed points Restatement loops are copies held by the context window: a statement still in context is near-copied 3.4× [2.6, 4.5] ...	45%	0.85	2.3	147
H70	The semantic information of the artifact store The artifact store tells an agent where to work after a context erasure (89% return to own repo; 0.14 bits) but adds ...	50%	0.70	1.8	149
H71	Memory is a first-order homeostat Agent memory size reverts to an agent-specific set point on the consolidation clock (ϕ^+ 0.67 per cycle, 37/37 units) ...	55%	0.74	1.1	151
H72	Trap aging is input starvation Trap aging is not input starvation: the idle-gate age slope survives the starvation control in 4/4 powered periods (G...	60%	0.78	1.8	153

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H73	Style = weights + context + register Style is mostly an agent constant: three components explain 0.72 of systematic style variance (34/34 periods), the ag...	45%	0.80	1.6	155
H74	An operator-grade change detector A schema diff of the raw record is an operator-grade change detector: 0/34 placebo false alarms; it dates NE14, NE17 ...	45%	0.63	2.0	157
H75	A thermodynamic speed limit on re-allocation A kickoff that names its target moves the swarm at the Wasserstein speed limit (slack S = 1.0–1.4, settling 0.25–0.5 ...	40%	0.47	1.4	159
H76	Excess vs housekeeping irreversibility The swarm’s excess behavioral irreversibility is the scheduler’s day end: 80–100% of debiased excess lies in the firs...	40%	0.73	1.5	161
H77	Repos as replicators: the selection-resolution bound Repo recruitment flux (σ^*) in herding and own-artifact weeks stays inside the band neutral copying produces on the sa...	40%	0.55	0.7	163
H78	Replicator growth order: herding vs division of labor Replicator growth order p is not identifiable from commit logs: first-order copying plus fitness spread reproduces ev...	40%	0.55	0.7	165
H79	An artifact-only autocatalytic set No artifact-only autocatalytic set exists beyond first order: self-catalysis (cron streams, repos running their own c...	40%	0.72	1.3	167
H80	Assembly index vs compression as an agent signature The assembly index adds nothing over compression for telling agent work from scripted work (ΔAUC -0.007 to +0.001, po...	40%	0.68	1.2	169
H81	Culture beyond composition: the emergent slow mode In regime I, village content has a collective slow mode beyond its members’ own vectors and the measured goal fields:...	65%	0.61	1.1	171
H82	Remanence: the endogenous part of the mean field Day-1 content of a new goal carries the previous period’s village content ($\Delta\gamma_1 \approx 0.15$, 3× the null, 25 boundaries, bo...	45%	0.73	1.1	173
H83	Enculturation of newcomers Newcomers match the village’s day-to-day content from tenure day 2 (gap to the veteran centroid -0.020 [-0.056, +0.02...	40%	0.53	0.5	175
H84	History-search outage as a collective-memory scramble The 03-31/04-01 history-search outage cost no measurable continuity or commits: a dip of 0.30 at the mean searcher do...	40%	0.42	0.4	177

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H85	Scaling of outputs with N Across non-holdout period units with regime intercepts, room message volume per scheduled hour grows sublinearly with...	45%	0.64	1.6	179
H86	Taylor’s law as a shared-field gauge Taylor’s c and exponent b are not field gauges for clocked agents ($b \approx 0.85 \pm 1$); the pair-covariance coefficient $c_{\dots}$	45%	0.68	1.5	181
H87	κ table: commits per bit by channel Among the swarm’s natural information channels, only the context window carries identifiable output value per bit ($\kappa_{\dots}$	45%	0.63	1.6	183
H88	Collective memory decays biexponentially Attention to a finished goal period’s terms falls within about a week ($\tau \sim 5.5$ village days) to a small floor (1.9% of ...	42%	0.62	0.9	185
H89	Price equation for village culture Within a goal period, village content, style and conventions change by transmission among stayers (shares 0.95–0.98, ...	45%	0.78	1.0	187
H90	Collective entropy production beyond the parts The swarm has one measurable collective arrow of time: in G51 talk, agents start talking in the minute after agents w...	45%	0.41	0.6	189
H91	Eigenvector rotation as a reorganization signal At non-holdout goal kickoffs at real counts (median $N \sim 10$, $W \sim 8$), day-to-day eigenvector rotation of the content matr...	45%	0.69	0.7	191
H92	RMT-cleaned matrices forecast tomorrow’s alignment Clipping eigenvalues at a calibrated edge forecasts tomorrow’s agent content correlations better than the raw matrix ...	45%	0.71	1.1	193
H93	Project choice as a Brock–Durlauf equilibrium Project choice is not bistable: every period sits below its Brock–Durlauf multiplicity boundary ($\beta J / \gamma_c \leq 0.86$ at poi...	40%	0.56	0.6	195
H94	Work allocation as a max-entropy equilibrium Work episodes are allocated at maximum entropy given agent activity, repo sizes, ownership and rooms: the residual ab...	40%	0.61	0.9	197
H95	Speed-limit slack as a kickoff-specificity gauge Kickoff goal type sets re-allocation speed: goals that assign a concrete artifact settle at slack 1.0–2.5 within 0.25...	40%	0.35	1.0	199
H96	Goal switches show hysteresis: a coercive field A goal switch is a quench, not a hysteresis loop: day-1 persistence of the old-state direction is R_1 0.18–0.27 vs 0.8...	45%	0.74	1.1	201

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H97	A restoring-force law for the kickoff quench A kickoff erases agents' pre-kickoff content positions beyond ordinary days (extra forgetting $\Delta\rho +0.26 \pm 0.04$, positi...	45%	0.70	1.1	203
H98	#51 is a random-field system #51's static content structure is a random field with a weak pull: after style removal the disorder ratio is R 0.67–0...	40%	0.49	0.5	205
H99	Fluctuations predict relaxation (mean-field Glauber) In regime III the collective talk memory beyond mean-field Glauber is read-out coupling, not the scheduler: it surviv...	65%	0.63	0.9	207
H100	#best/#rest: explicit vs spontaneous symmetry breaking In regime-III two-room periods, the #best/#rest content separation exceeds random regroupings after removing members'...	45%	0.71	1.1	209
H101	Pairwise Ising vs full multi-information in co-usage Convention co-usage is explained by pairwise couplings plus shared fields (field-removed ρ_F median 0.91; higher-orde...	45%	0.42	0.4	211
H102	Domain walls between rooms, set by bridging agents In #51g (the #focus weeks), room hoppers are not intermediate on the inter-domain axis and hopping read-outs do not m...	45%	0.51	0.8	213
H103	Nights demagnetize: remanence decays per night Nights do not demagnetize content: kickoff remanence decays on the active-hour clock, not per night (nights win 3–5 o...	45%	0.59	0.6	215
H104	Barkhausen avalanches when a field is stepped Human messages do not release switch avalanches: the #51 step response is $\times 1.07$ – 1.09 (upper bounds $\times 1.18$ – 1.32 , below ...	40%	0.49	0.7	217
H105	Goals act on a two-state order parameter On-goal occupancy rises from the decoy floor (≈ 0.03) to 0.23–0.44 after a kickoff and its variance grows as $p(1 - p)$...	40%	0.36	0.2	219
H106	A finite magnet's magnetization wanders as $1/N$: the slow mode's decay... In regime I (24 goals, N 4–12) the slow mode's finite-size exponent is not identifiable: $\alpha_k +0.40$ [-1.62, 2.41] bge ...	40%	0.62	0.2	221
H107	Does the room split grow from zero (an instability) or appear at once... Room content splits do not have one onset shape: in #37 and #39 the split starts near zero and grows over the week (r...	45%	0.35	0.2	223
H108	Goldstone wandering: the spontaneous room direction should drift, whi... Within a goal period the room-difference content direction persists day to day (noise-corrected P(1) 0.69–0.95 in 6/8...	45%	0.42!	0.2	225

ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H109	The forced erasure is a demagnetizing pulse: where is the room's orde... A forced context erasure does not drop an agent's alignment with its room's spontaneous axis (δ_F -0.08 [-0.23, 0.07]...	45%	0.62	0.9	227
H110	Exchange bias: an agent's own artifact shifts its goal-switch loop An agent's own live repo does not pin its old goal state: a small day-1 residual (dR 0.30 [-0.01, 0.57] bge) is not t...	40%	0.35	0.2	229
H111	A fluctuation-response sum rule for talk: Fano factor = $1/(1 - g)^2$ In regime III, talk obeys a parameter-free fluctuation-response sum rule: the collective Fano ratio of 10-min talk co...	50%	0.65	1.3	231
H112	Crossing claims anti-coordinate: parallel updates make 2-cycles Agents who join the same project do not form 2-cycles: once chat engagement is matched, departures do not depend on w...	40%	0.44	0.2	233
H113	Read-out channel capacity: information per call grows as $k^{\wedge}(1-\beta) \approx k^{\wedge}...$ The read-out channel has an intermediate capacity: content uptake per call grows as $k^{\wedge}0.31$ [0.25, 0.37] on #51 timer-...	45%	0.64	1.0	235
H114	A Griffiths phase: rare strong pairs give a subcritical swarm heavy-t... Reply chains are heavier than a uniform subcritical process, but thread momentum explains it, not rare strong pairs: ...	45%	0.54	0.5	237
H115	Blind role recovery in the debate week: the judge is a sink of antisy... In #12, per-call talk coupling runs across debate teams, not within them (J_ST - J_OT = -0.16 $\pm$ 0.08, p 0.01; J_OT 0....	40%	0.50	0.5	239
H116	The election as a coupling step: influence should flow toward the win... Winning the #26 election did not raise the per-call talk response to the winner (+0.09 $\pm$ 0.33; a step $\geq$ 0.75 Ising un...	40%	0.42	0.4	241
H117	J should not move when the rules do: the mid-goal reset in #32 (NE35) At the one step whose talk fields clearly moved (the #41 $\rightarrow$ #42 goal switch), the kinetic-Ising coupling matrix stayed...	45%	0.52	0.8	243
H118	The forked RPG is two replicas of one dynamics: damage spreading betw... After the RPG fork (NE15), the two rooms' content separates within the first active hour (overlap 0.51 vs 0.80 for ra...	40%	0.52	0.5	245
H119	The room merge switches a known adjacency on and off: recover it from... Merging the rooms (NE42) did not raise cross-group talk coupling (-0.02 [-0.13, 0.09], power 0.95 at J 0.3), and with...	45%	0.54	0.8	247
H120	Is a goal period a stationary NESS? Within-period drift of J and EP (... Goal periods longer than about a week are not stationary: #38 drifts in both activity and talk couplings (half-vs-hal...	45%	0.68	1.7	249
ID	Hypothesis (scored claim)	Complete	Credence	EU	p.
H123	Regime I's turn order is a sweep: equilibrium-looking statistics with... Regime I's update order is self-clocked asynchronous, not a sweep (41/41 units; order predictability median 0.04 vs $\geq$...	45%	0.67	0.7	251
H124	Four agents for weeks: an exact kinetic Ising benchmark for the mean-... Mean-field inversions (TAP, Sessak-Monasson, naive) match exact kinetic-Ising ML only when agents' self-coupling is w...	45%	0.59	0.9	253
H128	Free kickoffs coarsen, named kickoffs freeze: a kinetic Potts quench After a goal kickoff, the domain-wall fraction of agents over projects does not separate free from named kickoffs (on...	45%	0.51	0.5	255
H129	Project hopping carries cycle currents: the sticky Potts walker break... Project hopping carries no circulation beyond the drift from older to newer projects (rotation-summed cycle affinity ...	45%	0.52	0.5	257
H130	#51 agents are Ornstein-Uhlenbeck particles in private wells: reads k... In #51 a read pulls the reader's next statement toward the sender beyond the in-flight placebo (J_K 0.044 [0.038, 0.0...	50%	0.64	1.0	259
H131	Assigned antagonism switches off in one read-out after the prize goes... Assigned antagonism is fully off by each rival's first reply after it reads the verdict (#12: 0/21 first post-read ri...	50%	—	—	261
H132	Debate is a two-sublattice limit cycle: topic ping-pong at lag one Debate content does not alternate between teams turn by turn (Δ +0.05, inside the turn-shuffled null; power 1.00 at c...	45%	0.44	0.2	263